


```
My dataset is available in googledrive; so I am accessing my d
from google.colab import drive
drive.mount('/content/drive', force_remount=True)
Once this is executed, you will see your drive appearing
```



```
Mounted at /content/drive
```

```
# Define paths to the dataset
train_dir = '/content/drive/MyDrive/horse-or-human/test'
test_dir = '/content/drive/MyDrive/horse-or-human/train'
```



```
import os
print(len(os.listdir('/content/drive/MyDrive/horse-or-human/'))
print(len(os.listdir('/content/drive/MyDrive/horse-or-human/'))
print(len(os.listdir('/content/drive/MyDrive/horse-or-human/'))
print(len(os.listdir('/content/drive/MyDrive/horse-or-human/'))
```



```
500
527
500
527
```

```
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator

# Create ImageDataGenerator for training set
train_datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=0.2 # Split 20% of the images for validation
)

# Load and prepare training data
train_data = train_datagen.flow_from_directory(
    train_dir,
    target_size=(256,256),
    batch_size=32,
    class_mode='binary', # 'binary' for binary classification
    subset='training' # Specify 'training' for the training
)

# Create ImageDataGenerator for validation set
validation_datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=0.2 # Note: Using the same validation split
)

# Load and prepare validation data
validation_data = validation_datagen.flow_from_directory(
    train_dir,
    target_size=(256,256),
    batch_size=32,
    class_mode='binary',
    subset='validation' # Specify 'validation' for the validation
)
```

Found 822 images belonging to 2 classes.  
Found 205 images belonging to 2 classes.

```
from tensorflow import keras
from keras import Sequential
from keras.layers import Dense,Conv2D,MaxPooling2D,Flatten,Batch
```

```
# create CNN model - custom-made

model = Sequential()

model.add(Conv2D(32,kernel_size=(3,3),padding='valid',activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

model.add(Conv2D(64,kernel_size=(3,3),padding='valid',activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

model.add(Conv2D(128,kernel_size=(3,3),padding='valid',activation='relu'))
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid'))

model.add(Flatten())

model.add(Dense(128,activation='relu')) #feature reduction
model.add(Dense(64,activation='relu'))
model.add(Dense(1,activation='sigmoid')) #output layer
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional.py:105: UserWarning: This class is deprecated. It is recommended to use the Keras API instead. See: https://keras.io/
```


```
from keras.optimizers import Adam
model.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy')
```

```
history = model.fit(train_data, epochs=20, validation_data=validation_data)


/usr/local/lib/python3.11/dist-packages/keras/src/trainers/d
    self._warn_if_super_not_called()
Epoch 1/20
26/26 193s 7s/step - accuracy: 0.5979 -
Epoch 2/20
26/26 138s 5s/step - accuracy: 0.9563 -
Epoch 3/20
26/26 140s 5s/step - accuracy: 0.9933 -
Epoch 4/20
26/26 139s 5s/step - accuracy: 0.9959 -
Epoch 5/20
26/26 149s 6s/step - accuracy: 0.9954 -
Epoch 6/20
26/26 139s 5s/step - accuracy: 0.9374 -
Epoch 7/20
26/26 138s 5s/step - accuracy: 0.9776 -
Epoch 8/20
26/26 148s 6s/step - accuracy: 0.9873 -
Epoch 9/20
26/26 149s 6s/step - accuracy: 1.0000 -
Epoch 10/20
26/26 149s 6s/step - accuracy: 1.0000 -
Epoch 11/20
26/26 209s 6s/step - accuracy: 1.0000 -
Epoch 12/20
26/26 139s 5s/step - accuracy: 1.0000 -
Epoch 13/20
26/26 137s 5s/step - accuracy: 1.0000 -
Epoch 14/20
26/26 137s 5s/step - accuracy: 1.0000 -
Epoch 15/20
26/26 137s 5s/step - accuracy: 1.0000 -
Epoch 16/20
26/26 139s 5s/step - accuracy: 1.0000 -
Epoch 17/20
26/26 150s 6s/step - accuracy: 1.0000 -
Epoch 18/20
26/26 150s 6s/step - accuracy: 1.0000 -
Epoch 19/20
26/26 190s 5s/step - accuracy: 1.0000 -
Epoch 20/20
26/26 152s 6s/step - accuracy: 1.0000 -
```

```
import matplotlib.pyplot as plt

plt.plot(history.history['accuracy'],color='red',label='train accuracy')
plt.plot(history.history['val_accuracy'],color='blue',label='validation accuracy')
plt.legend()
plt.show()
```



```
plt.plot(history.history['loss'],color='red',label='train')
plt.plot(history.history['val_loss'],color='blue',label='val')
plt.legend()
plt.show()
```



```
test_datagen = ImageDataGenerator(rescale=1./255)
test_data = test_datagen.flow_from_directory(
    test_dir,
    target_size=(256,256),
    batch_size=32,
    class_mode='binary'
)
```

Found 1027 images belonging to 2 classes.

```
#predict the test data
predictions = model.predict(test_data)
```

**33/33** ————— **118s** 4s/step

```
from sklearn.metrics import confusion_matrix, classification_report

# Assuming you have ground truth labels (true_labels) and predictions
true_labels = test_data.classes
predicted_labels = (predictions > 0.5).astype(int) # Adjust threshold

# Calculate confusion matrix
cm = confusion_matrix(true_labels, predicted_labels)

# Print confusion matrix
print("Confusion Matrix:")
print(cm)

# Print classification report
print("Classification Report:")
print(classification_report(true_labels, predicted_labels))
```

Confusion Matrix:

```
[[247 253]
 [266 261]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.48	0.49	0.49	500
1	0.51	0.50	0.50	527
accuracy			0.49	1027
macro avg	0.49	0.49	0.49	1027
weighted avg	0.49	0.49	0.49	1027

```
# create CNN model

model = Sequential()

model.add(Conv2D(32,kernel_size=(3,3),padding='valid',activation='relu')) # added to reduce overfitting
model.add(BatchNormalization()) # added to reduce overfitting
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid')) # added to reduce overfitting

model.add(Conv2D(64,kernel_size=(3,3),padding='valid',activation='relu')) # added to reduce overfitting
model.add(BatchNormalization()) # added to reduce overfitting
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid')) # added to reduce overfitting

model.add(Conv2D(128,kernel_size=(3,3),padding='valid',activation='relu')) # added to reduce overfitting
model.add(BatchNormalization()) # added to reduce overfitting
model.add(MaxPooling2D(pool_size=(2,2),strides=2,padding='valid')) # added to reduce overfitting

model.add(Flatten())

model.add(Dense(128,activation='relu')) #feature reduction
model.add(Dropout(0.1)) # added to reduce overfitting
model.add(Dense(64,activation='relu'))
model.add(Dropout(0.1)) # added to reduce overfitting
model.add(Dense(1,activation='sigmoid')) #output layer

/usr/local/lib/python3.11/dist-packages/keras/src/layers/convolutional.py:105: UserWarning: This class is deprecated. Use Conv2D instead.
  super().__init__(activity_regularizer=activity_regularizer,
```

```
from keras.optimizers import Adam
model.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy')
```

```
history = model.fit(train_data, epochs=20, validation_data=validation_data)
```

```
Epoch 1/20
26/26 235s 9s/step - accuracy: 0.8428 -
Epoch 2/20
26/26 209s 8s/step - accuracy: 0.9778 -
Epoch 3/20
26/26 208s 8s/step - accuracy: 0.9895 -
Epoch 4/20
26/26 216s 8s/step - accuracy: 0.9832 -
Epoch 5/20
26/26 261s 8s/step - accuracy: 0.9872 -
Epoch 6/20
26/26 216s 8s/step - accuracy: 0.9957 -
Epoch 7/20
26/26 209s 8s/step - accuracy: 0.9931 -
Epoch 8/20
26/26 207s 8s/step - accuracy: 0.9929 -
Epoch 9/20
26/26 261s 8s/step - accuracy: 0.9981 -
Epoch 10/20
26/26 207s 8s/step - accuracy: 0.9972 -
Epoch 11/20
26/26 208s 8s/step - accuracy: 0.9887 -
Epoch 12/20
26/26 260s 8s/step - accuracy: 0.9919 -
Epoch 13/20
26/26 207s 8s/step - accuracy: 0.9850 -
Epoch 14/20
26/26 207s 8s/step - accuracy: 0.9980 -
Epoch 15/20
26/26 209s 8s/step - accuracy: 0.9970 -
Epoch 16/20
26/26 206s 8s/step - accuracy: 0.9986 -
Epoch 17/20
26/26 215s 8s/step - accuracy: 0.9976 -
Epoch 18/20
26/26 206s 8s/step - accuracy: 0.9953 -
Epoch 19/20
26/26 214s 8s/step - accuracy: 0.9968 -
Epoch 20/20
26/26 206s 8s/step - accuracy: 0.9998 -
```

```
test_data = test_datagen.flow_from_directory(  
    test_dir,  
    target_size=(256,256),  
    batch_size=32,  
    class_mode='binary'  
)
```

```
Found 1027 images belonging to 2 classes.
```

```
predictions = model.predict(test_data)
```

```
/usr/local/lib/python3.11/dist-packages/keras/src/trainers/da  
    self._warn_if_super_not_called()  
33/33 ━━━━━━━━━━━━━━━━ 58s 2s/step
```

```
from sklearn.metrics import confusion_matrix, classification_report

# Assuming you have ground truth labels (true_labels) and predictions (predictions)
true_labels = test_data.classes
predicted_labels = (predictions > 0.5).astype(int) # Adjust threshold

# Calculate confusion matrix
cm = confusion_matrix(true_labels, predicted_labels)

# Print confusion matrix
print("Confusion Matrix:")
print(cm)

# Print classification report
print("Classification Report after applying techniques to handle ove")
print(classification_report(true_labels, predicted_labels))
```

Confusion Matrix:

```
[[263 237]
 [277 250]]
```

Classification Report after applying techniques to handle ove

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.49	0.53	0.51	500
1	0.51	0.47	0.49	527

accuracy			0.50	1027
macro avg	0.50	0.50	0.50	1027
weighted avg	0.50	0.50	0.50	1027

```
import tensorflow as tf
from tensorflow import keras
from keras import Sequential
from keras.layers import Dense, Conv2D, MaxPooling2D, Flatten

resnet_model = Sequential()      #ResNet50 is a pre-trained model
pretrained_model = tf.keras.applications.ResNet50(include_top=False,
                                                    input_shape=(224, 224, 3),
                                                    pooling='max',
                                                    weights='imagenet')

for layer in pretrained_model.layers:
    layer.trainable = False      #all layers in the Resnet50 is frozen

resnet_model.add(pretrained_model)
resnet_model.add(Flatten())
resnet_model.add(Dense(512, activation = 'relu'))
resnet_model.add(Dense(1, activation = 'sigmoid'))
```

```
Downloading data from https://storage.googleapis.com/tensorflow-1.x-20180228-144212-094765736/94765736 0s 0us/step
```

```
from keras.optimizers import Adam
resnet_model.compile(optimizer=Adam(learning_rate=0.001), loss='binary_crossentropy', metrics=['accuracy'])
```

```
history = resnet_model.fit(train_data, epochs=20, validation_
```

```
Epoch 1/20
26/26 268s 10s/step - accuracy: 0.5358 -
Epoch 2/20
26/26 255s 10s/step - accuracy: 0.7011 -
Epoch 3/20
26/26 286s 11s/step - accuracy: 0.8359 -
Epoch 4/20
26/26 290s 10s/step - accuracy: 0.9046 -
Epoch 5/20
26/26 254s 10s/step - accuracy: 0.8924 -
Epoch 6/20
26/26 287s 11s/step - accuracy: 0.9405 -
Epoch 7/20
26/26 255s 10s/step - accuracy: 0.9153 -
Epoch 8/20
26/26 255s 10s/step - accuracy: 0.9704 -
Epoch 9/20
26/26 255s 10s/step - accuracy: 0.9342 -
Epoch 10/20
26/26 254s 10s/step - accuracy: 0.9681 -
Epoch 11/20
26/26 255s 10s/step - accuracy: 0.9707 -
Epoch 12/20
26/26 286s 11s/step - accuracy: 0.9557 -
Epoch 13/20
26/26 255s 10s/step - accuracy: 0.9535 -
Epoch 14/20
26/26 256s 10s/step - accuracy: 0.9588 -
Epoch 15/20
26/26 256s 10s/step - accuracy: 0.9565 -
Epoch 16/20
26/26 256s 10s/step - accuracy: 0.9788 -
Epoch 17/20
26/26 292s 11s/step - accuracy: 0.9793 -
Epoch 18/20
26/26 254s 10s/step - accuracy: 0.9511 -
Epoch 19/20
26/26 254s 10s/step - accuracy: 0.9780 -
Epoch 20/20
26/26 256s 10s/step - accuracy: 0.9700 -
```

```
predictions = resnet_model.predict(test_data)
```

```
33/33 265s 8s/step
```

```
from sklearn.metrics import confusion_matrix, classification_report

# Assuming you have ground truth labels (true_labels) and predictions (predictions)
true_labels = test_data.classes
predicted_labels = (predictions > 0.5).astype(int) # Adjust threshold

# Calculate confusion matrix
cm = confusion_matrix(true_labels, predicted_labels)

# Print confusion matrix
print("Confusion Matrix:")
print(cm)

# Print classification report
print("Classification Report after applying techniques to handle overfitting")
print(classification_report(true_labels, predicted_labels))
```

Confusion Matrix:

```
[[259 241]
 [312 215]]
```

Classification Report after applying techniques to handle overfitting

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.45	0.52	0.48	500
1	0.47	0.41	0.44	527

accuracy			0.46	1027
macro avg	0.46	0.46	0.46	1027
weighted avg	0.46	0.46	0.46	1027

